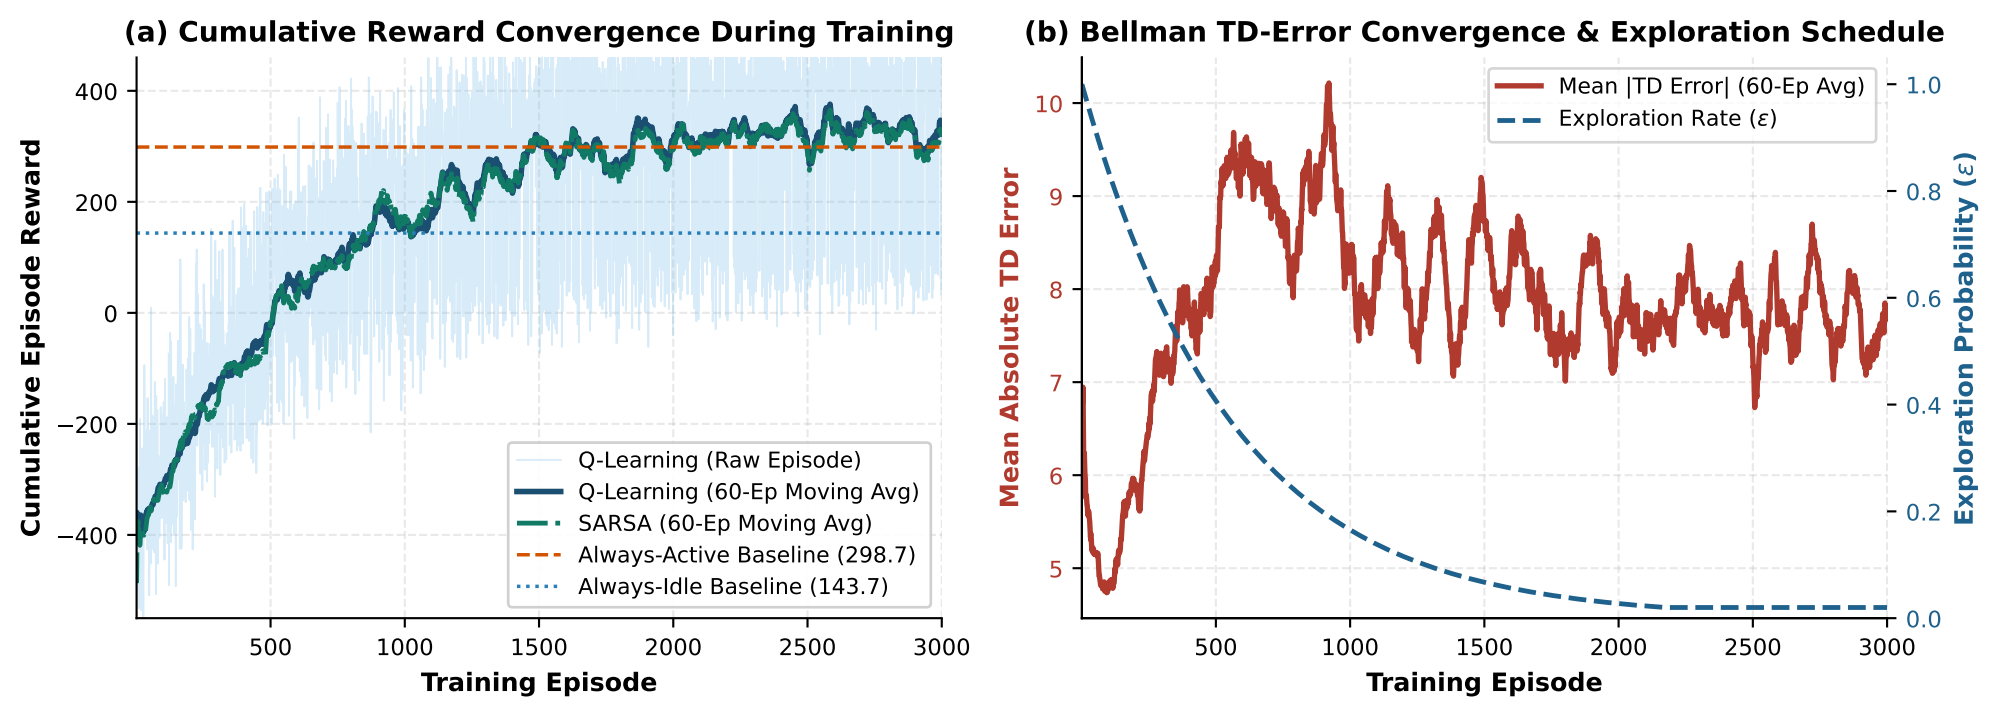

(a) Cumulative Reward Convergence During Training
Cumulative Episode Reward
Training Episode
Q-Learning (Raw Episode)
Q-Learning (60-Ep Moving Avg)
SARSA (60-Ep Moving Avg)
Always-Active Baseline (298.7)
Always-Idle Baseline (143.7)
(b) Bellman TD-Error Convergence & Exploration Schedule
Mean Absolute TD Error
Exploration Probability (ε)
Training Episode
Mean |TD Error| (60-Ep Avg)
Exploration Rate (ε)